



# Apache Kafka:

Powering Real-Time Data Pipelines and  
Event-Driven Architectures

# Agenda



01. Introduction

03. **Development  
and need for Kafka**

05. **Demo**

02. **Historical  
Background**

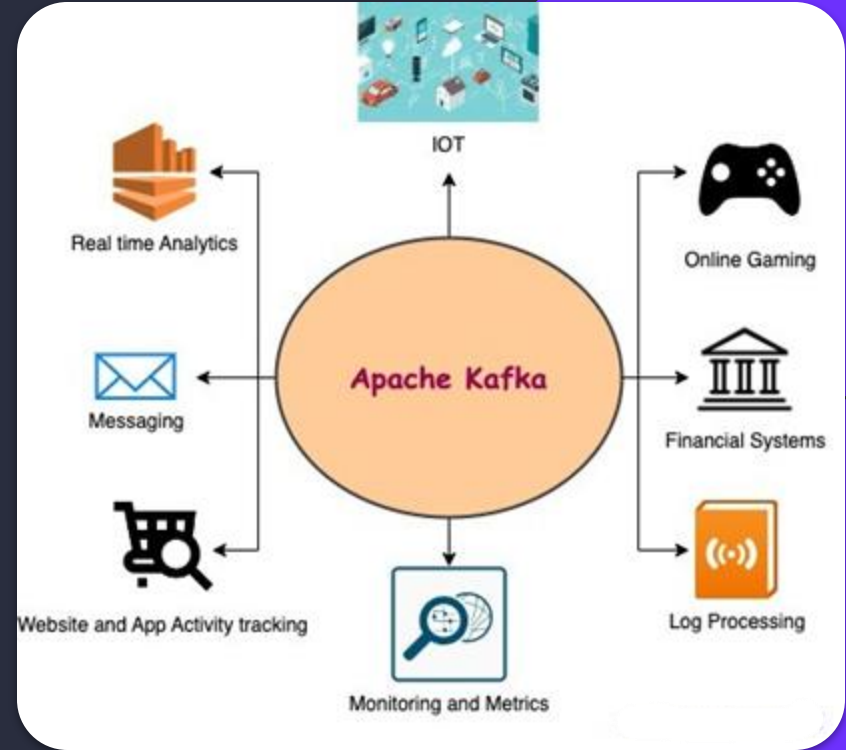
04. **Architecture**

# Kafka

The topic of our term paper is Apache Kafka. It is an open-source distributed event streaming platform used by thousands of companies for high-performance data pipelines, streaming analytics, data integration, and mission-critical applications

Examples:

- Ingesting clickstream data from websites into storage (e.g., S3, Snowflake, BigQuery).
- Streaming IoT sensor data into analytics systems.
- Sending data from microservices to analytics dashboards.



# Kafka: Who, When and Why?

Kafka was originally **created at LinkedIn in 2010** by engineers Jay Kreps, Neha Narkhede, and Jun Rao, who were frustrated by the limitations of traditional messaging systems like ActiveMQ and RabbitMQ in handling the massive scale of LinkedIn's data pipelines.

ActiveMQ and RabbitMQ could handle thousands of messages per second, but LinkedIn needed to handle millions. These systems stored messages in memory or relational databases, causing severe slowdowns under high load. Kafka, in contrast, writes messages sequentially to disk and batches network I/O, allowing near-file-system-level performance.

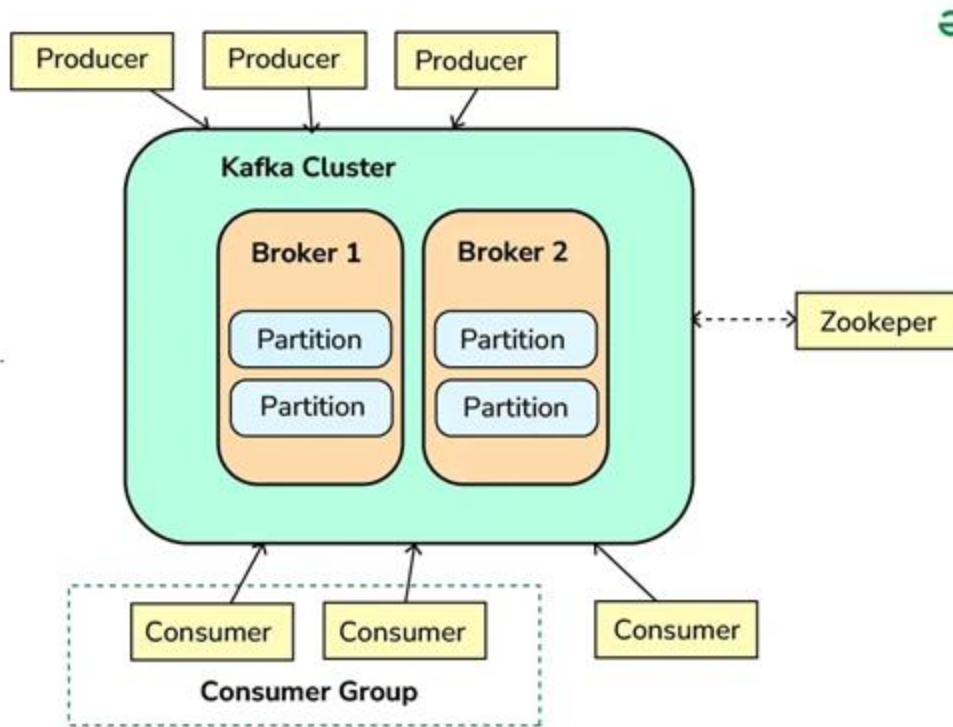
Their goal was to design a platform capable of ingesting billions of events per day, ranging from user activity logs and clickstream data to recommendation engine signals, while providing strong guarantees around scalability, durability, and fault tolerance. **This internal project was open-sourced under the Apache Software Foundation in 2011.**



# Kafka: Who, When and Why?

Feature	Traditional Messaging	Kafka
Scalability	Single broker or small cluster	Horizontally scalable across servers
Storage	Deletes after consumption	Persistent, replayable log
Speed	Thousands/sec	Millions/sec
Use cases	Notifications, transactions	Real-time data pipelines, analytics, ML
Integration	One-to-one	One-to-many

# Kafka Architecture



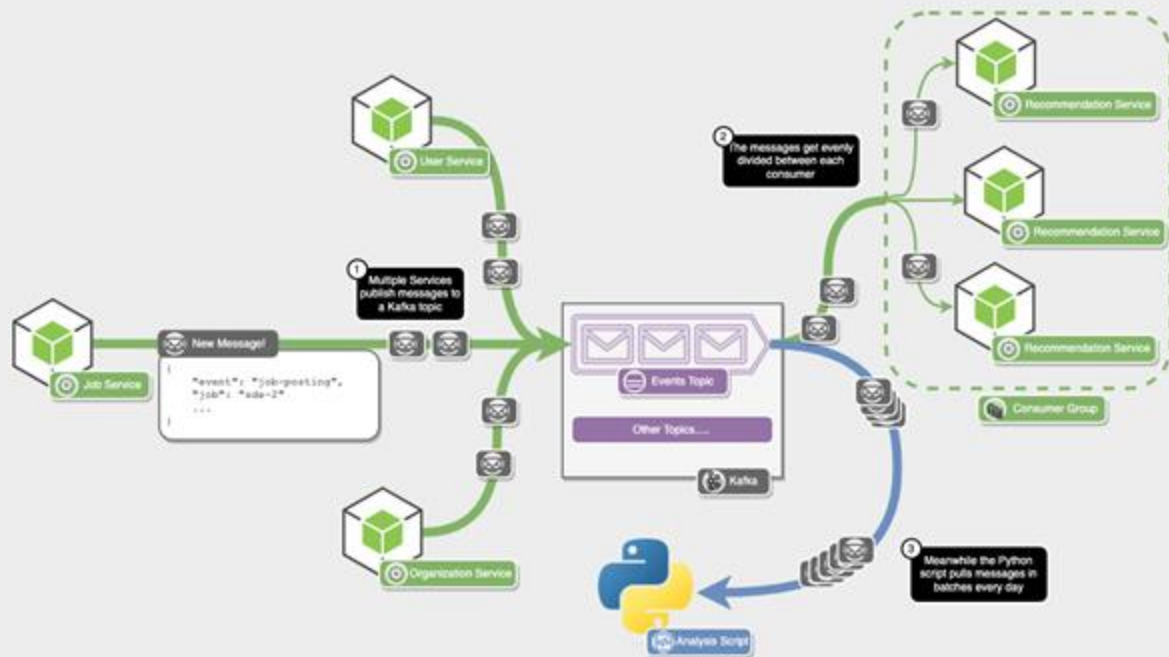
- **Producers** send data (events) into the **Kafka Cluster**.
- The **cluster** is made up of multiple **brokers** (servers) that store data.
- Each broker holds **partitions**, which keep messages in order for fast access.
- **Consumers** read data from these partitions, often as part of a **consumer group** to share the load.
- **ZooKeeper** (or **KRaft** in newer versions) manages coordination and health of brokers.

Together, this design makes Kafka **scalable, fault-tolerant, and real-time**.

Multiple services like the Job, User, and Organization services generate events (such as job postings) and **publish them to a Kafka topic**. Kafka acts as a **central event hub**, storing and streaming these messages in real time.

A **consumer group** (the Recommendation Service instances) subscribes to this topic, and Kafka **evenly distributes messages** across them for parallel processing.

Meanwhile, other consumers, like a **Python analysis script**, can independently **pull and process the same data later** — showing Kafka's ability to support both real-time and batch processing simultaneously.



# Kafka Techniques & Approaches

## 1. Distributed Publish–Subscribe System

Kafka enables multiple producers and consumers to exchange data in real time through topics — allowing systems to remain loosely coupled and highly scalable.

## 2. Partitioning for Parallelism

Topics are divided into partitions, letting multiple brokers handle data simultaneously. This ensures **high throughput** and **load balancing** across servers.

## 3. Replication for Reliability

Each partition has replicas across brokers to prevent data loss and support **fault-tolerant recovery** in case of hardware or network failures.

## 4. Stream Processing Capability

Kafka Streams API supports filtering, aggregating, and transforming messages in real time, turning raw events into actionable insights.

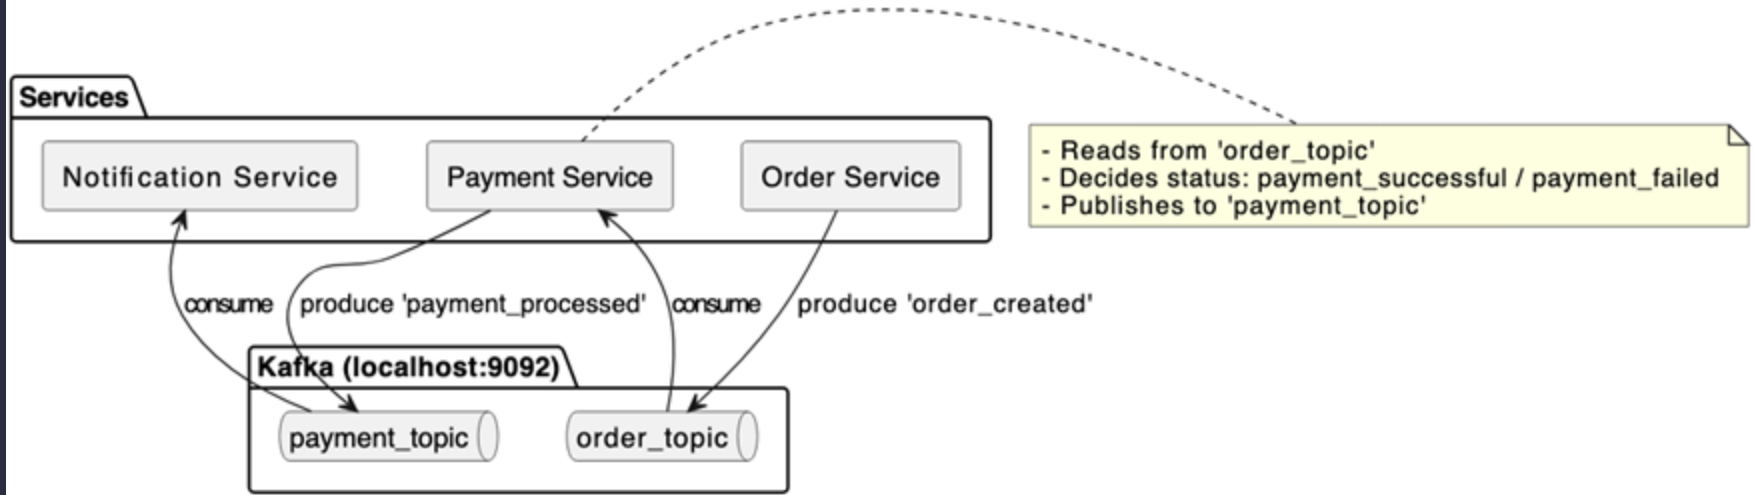
## 5. Seamless Integration

Kafka connects easily with **Spark, Flink, and Hadoop ecosystems**, making it central to **real-time analytics and event-driven architectures**.



# Demo Architecture

Kafka Microservices — Component View (as-coded)



# Sample Output

```
aryaman@Aryamans-MacBook-Air kafka_microservices_demo % docker compose up -d
[+] Running 2/2
 ✓ Container kafka_microservices_demo-zookeeper-1 Running
 ✓ Container kafka_microservices_demo-kafka-1 Started
aryaman@Aryamans-MacBook-Air kafka_microservices_demo % docker ps
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS
847a8141337a   confluentinc/cp-kafka:7.6.1        "/etc/confluent/dock_"  8 days ago    Up 6 seconds
_demo-kafka-1
315ec646bafc   confluentinc/cp-zookeeper:7.6.1    "/etc/confluent/dock_"  8 days ago    Up 3 minutes
_demo-zookeeper-1
aryaman@Aryamans-MacBook-Air kafka_microservices_demo % python order_service.py
Order Service started. Generating new orders...
Order Created: {'order_id': 1000, 'user_id': 73, 'amount': 453.91, 'status': 'created'}
Order Created: {'order_id': 1001, 'user_id': 43, 'amount': 41.58, 'status': 'created'}
Order Created: {'order_id': 1002, 'user_id': 76, 'amount': 412.97, 'status': 'created'}
Order Created: {'order_id': 1003, 'user_id': 40, 'amount': 219.16, 'status': 'created'}
Order Created: {'order_id': 1004, 'user_id': 84, 'amount': 111.55, 'status': 'created'}
Order Created: {'order_id': 1005, 'user_id': 25, 'amount': 427.34, 'status': 'created'}
Order Created: {'order_id': 1006, 'user_id': 99, 'amount': 390.53, 'status': 'created'}
Order Created: {'order_id': 1007, 'user_id': 87, 'amount': 309.26, 'status': 'created'}
Order Created: {'order_id': 1008, 'user_id': 1, 'amount': 159.22, 'status': 'created'}
Order Created: {'order_id': 1009, 'user_id': 61, 'amount': 17.69, 'status': 'created'}
Order Created: {'order_id': 1010, 'user_id': 95, 'amount': 50.73, 'status': 'created'}
Order Created: {'order_id': 1011, 'user_id': 86, 'amount': 295.61, 'status': 'created'}
Order Created: {'order_id': 1012, 'user_id': 6, 'amount': 398.75, 'status': 'created'}
Order Created: {'order_id': 1013, 'user_id': 81, 'amount': 115.58, 'status': 'created'}
Order Created: {'order_id': 1014, 'user_id': 82, 'amount': 322.48, 'status': 'created'}
Order Created: {'order_id': 1015, 'user_id': 81, 'amount': 22.5, 'status': 'created'}
Order Created: {'order_id': 1016, 'user_id': 63, 'amount': 459.19, 'status': 'created'}
```

# Sample Output

```
aryaman@Aryamans-MacBook-Air kafka_microservices_demo % docker compose up -d
[+] Running 2/2
 ✓ Container kafka_microservices_demo-zookeeper-1 Running
 ✓ Container kafka_microservices_demo-kafka-1 Started

aryaman@Aryamans-MacBook-Air kafka_microservices_demo % python payment_service.py
CONTAINER ID      847a8141337a
demo-kafka-1      315ec646bafc
demo-zookeeper

Payment Service started. Listening for new orders...
Received Order: {'order_id': 1000, 'user_id': 73, 'amount': 453.91, 'status': 'created'}
Payment Event Sent: {'order_id': 1000, 'user_id': 73, 'amount': 453.91, 'status': 'payment_failed'}

Received Order: {'order_id': 1001, 'user_id': 43, 'amount': 41.58, 'status': 'created'}
Payment Event Sent: {'order_id': 1001, 'user_id': 43, 'amount': 41.58, 'status': 'payment_successful'}

Received Order: {'order_id': 1002, 'user_id': 76, 'amount': 412.97, 'status': 'created'}
Payment Event Sent: {'order_id': 1002, 'user_id': 76, 'amount': 412.97, 'status': 'payment_failed'}

Received Order: {'order_id': 1003, 'user_id': 40, 'amount': 219.16, 'status': 'created'}
Payment Event Sent: {'order_id': 1003, 'user_id': 40, 'amount': 219.16, 'status': 'payment_successful'}

Received Order: {'order_id': 1004, 'user_id': 84, 'amount': 111.55, 'status': 'created'}
Payment Event Sent: {'order_id': 1004, 'user_id': 84, 'amount': 111.55, 'status': 'payment_successful'}

Received Order: {'order_id': 1005, 'user_id': 25, 'amount': 427.34, 'status': 'created'}
Payment Event Sent: {'order_id': 1005, 'user_id': 25, 'amount': 427.34, 'status': 'payment_successful'}

Received Order: {'order_id': 1006, 'user_id': 99, 'amount': 390.53, 'status': 'created'}
Payment Event Sent: {'order_id': 1006, 'user_id': 99, 'amount': 390.53, 'status': 'payment_failed'}

Received Order: {'order_id': 1007, 'user_id': 87, 'amount': 309.26, 'status': 'created'}
Payment Event Sent: {'order_id': 1007, 'user_id': 87, 'amount': 309.26, 'status': 'payment_failed'}
```

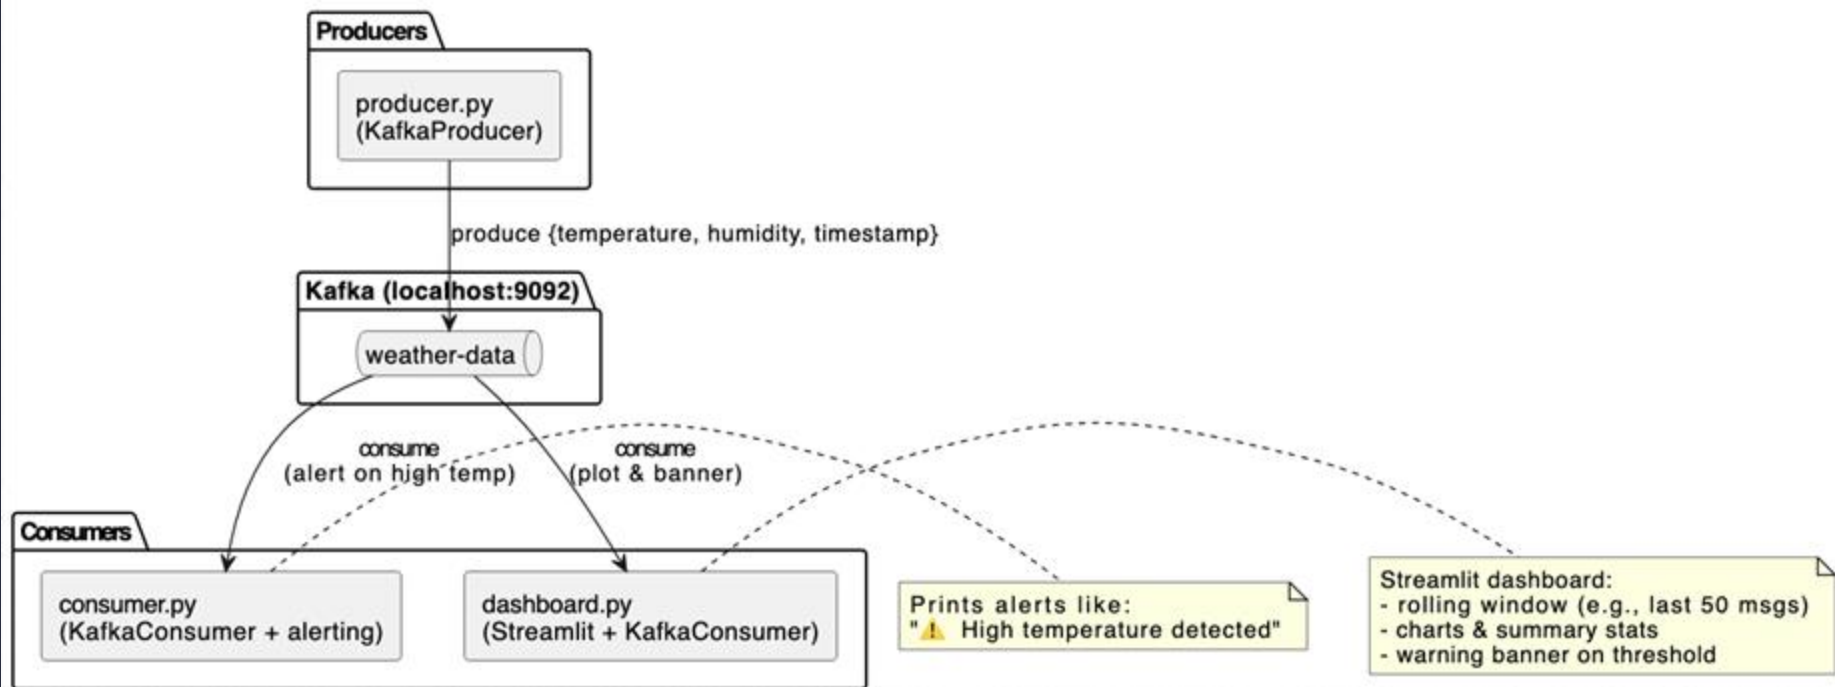
# Sample Output

```
aryaman@Aryamans-MacBook-Air kafka_microservices_demo % docker compose up -d
[+] Running 2/2
  ✓ Container kafka_microservices_demo-zookeeper-1 Running
  ✓ Container kafka_microservices_demo-kafka-1 Started

aryaman@Aryamans-MacBook-Air kafka_microservices_demo % python payment_service.py
CONTAINER ID      847a8141337a
demo-kafka-1      315ec646bafc
demo-zookeeper
Order Service
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Order Create
Received Order:
Payment Event Se
Payment Service started. Listening for new orders...
Received Order: {'order_id': 1000, 'user_id': 73, 'amount': 453.91, 'status': 'created'}
Payment Event Sent
Notification Service started. Waiting for payment updates...
Received Order:
Payment Event Se
Payment Failed for Order 1000 (User 73)
Email sent to customer: payment failed. Please retry.
Received Order:
Payment Event Se
Payment Successful for Order 1001 (User 43)
Email sent to customer confirming successful order.
Received Order:
Payment Event Se
Payment Failed for Order 1002 (User 76)
Email sent to customer: payment failed. Please retry.
Received Order:
Payment Event Se
Payment Successful for Order 1003 (User 40)
Email sent to customer confirming successful order.
Received Order:
Payment Event Se
Payment Successful for Order 1004 (User 84)
Email sent to customer confirming successful order.
Received Order:
Payment Event Se
Payment Successful for Order 1005 (User 25)
Email sent to customer confirming successful order.
Payment Failed for Order 1006 (User 99)
Email sent to customer: payment failed. Please retry.
```

# Demo Architecture

IoT Weather Stream — Component View (as-coded)





# Sample Output

```
🔥 Kafka Producer started... streaming temperature & humidity data
Produced: {'sensor_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}
Produced: {'sensor_id': 387, 'temperature': 31.33, 'humidity': 72.8, 'timestamp': '2025-10-30 19:02:34'}
Produced: {'sensor_id': 198, 'temperature': 35.67, 'humidity': 66.24, 'timestamp': '2025-10-30 19:02:35'}
Produced: {'sensor_id': 616, 'temperature': 32.38, 'humidity': 60.81, 'timestamp': '2025-10-30 19:02:36'}
Produced: {'sensor_id': 228, 'temperature': 28.62, 'humidity': 62.56, 'timestamp': '2025-10-30 19:02:37'}
Produced: {'sensor_id': 243, 'temperature': 23.51, 'humidity': 41.97, 'timestamp': '2025-10-30 19:02:38'}
Produced: {'sensor_id': 228, 'temperature': 21.27, 'humidity': 75.02, 'timestamp': '2025-10-30 19:02:39'}
Produced: {'sensor_id': 436, 'temperature': 23.1, 'humidity': 50.03, 'timestamp': '2025-10-30 19:02:40'}
Produced: {'sensor_id': 492, 'temperature': 37.15, 'humidity': 36.14, 'timestamp': '2025-10-30 19:02:41'}
Produced: {'sensor_id': 224, 'temperature': 32.47, 'humidity': 49.07, 'timestamp': '2025-10-30 19:02:42'}
Produced: {'sensor_id': 540, 'temperature': 38.71, 'humidity': 31.43, 'timestamp': '2025-10-30 19:02:43'}
Produced: {'sensor_id': 635, 'temperature': 25.99, 'humidity': 39.88, 'timestamp': '2025-10-30 19:02:44'}
Produced: {'sensor_id': 592, 'temperature': 21.73, 'humidity': 47.49, 'timestamp': '2025-10-30 19:02:45'}
Produced: {'sensor_id': 847, 'temperature': 23.04, 'humidity': 66.43, 'timestamp': '2025-10-30 19:02:46'}
Produced: {'sensor_id': 597, 'temperature': 30.83, 'humidity': 65.59, 'timestamp': '2025-10-30 19:02:47'}
Produced: {'sensor_id': 593, 'temperature': 34.57, 'humidity': 54.17, 'timestamp': '2025-10-30 19:02:48'}
Produced: {'sensor_id': 785, 'temperature': 39.26, 'humidity': 51.04, 'timestamp': '2025-10-30 19:02:49'}
Produced: {'sensor_id': 362, 'temperature': 39.89, 'humidity': 31.26, 'timestamp': '2025-10-30 19:02:50'}
Produced: {'sensor_id': 541, 'temperature': 22.99, 'humidity': 74.48, 'timestamp': '2025-10-30 19:02:51'}
Produced: {'sensor_id': 495, 'temperature': 39.63, 'humidity': 36.27, 'timestamp': '2025-10-30 19:02:52'}
Produced: {'sensor_id': 469, 'temperature': 37.1, 'humidity': 71.31, 'timestamp': '2025-10-30 19:02:53'}
Produced: {'sensor_id': 566, 'temperature': 39.72, 'humidity': 36.07, 'timestamp': '2025-10-30 19:02:54'}
Produced: {'sensor_id': 439, 'temperature': 27.7, 'humidity': 61.46, 'timestamp': '2025-10-30 19:02:55'}
Produced: {'sensor_id': 809, 'temperature': 38.51, 'humidity': 77.44, 'timestamp': '2025-10-30 19:02:56'}
Produced: {'sensor_id': 400, 'temperature': 35.45, 'humidity': 57.46, 'timestamp': '2025-10-30 19:02:57'}
Produced: {'sensor_id': 678, 'temperature': 31.51, 'humidity': 60.09, 'timestamp': '2025-10-30 19:02:58'}
```

# Sample Output

🔥 Kafka Producer started... streaming temperature & humidity data

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}


Produced: {'sensor\_id': 979, 'temperature': 34.84, 'humidity': 79.47, 'timestamp': '2025-10-30 19:02:33'}

## 🔥 Real-Time Weather Sensor Dashboard


⚠️ ALERT: High temperature detected (36.7°C)

sensor_id	temperature	humidity	timestamp
40	776	36.76	48.89 2025-10-30 19:04:00
46	822	30.84	79.44 2025-10-30 19:04:08
47	980	21.84	52.94 2025-10-30 19:04:06
46	329	31.6	34.67 2025-10-30 19:04:07
40	189	38.58	20.88 2025-10-30 19:04:06
44	405	39.1	47.22 2025-10-30 19:04:05
40	289	22.85	55.54 2025-10-30 19:04:04
42	888	36.37	42.28 2025-10-30 19:04:03
44	444	23.81	51.89 2025-10-30 19:04:02
40	629	30	58.14 2025-10-30 19:04:01

### 📊 Temperature (°C)



### 💧 Humidity (%)



**Thank you!**

